

# Time, Clocks, and the Ordering of Events in a Distributed System

a paper by Leslie Lamport

David Murray (@fernomac)

PWL @ Seattle, 2020-10-01

Operating  
Systems

R. Stockton Gaines  
Editor

---

# Time, Clocks, and the Ordering of Events in a Distributed System

Leslie Lamport  
Massachusetts Computer Associates, Inc.

Communications  
of  
the ACM

July 1978  
Volume 21  
Number 7

Network Working Group  
RFC # 677  
NIC # 31507

Paul R. Johnson (BBN-TENEX)  
Robert H. Thomas (BBN-TENEX)  
January 27, 1975

## The Maintenance of Duplicate Databases

### Preface:

This RFC is a working paper on the problem of maintaining duplicated databases in an ARPA-like network. It briefly discusses the general duplicate database problem, and then outlines in some detail a solution for a particular type of duplicate database. The concepts developed here were used in the design of the User Identification Database for the TIP user authentication and accounting system. We believe that these concepts are generally applicable to distributed database problems.

Operating  
Systems

R. Stockton Gaines  
Editor

---

# Time, Clocks, and the Ordering of Events in a Distributed System

Leslie Lamport  
Massachusetts Computer Associates, Inc.

Communications  
of  
the ACM


July 1978  
Volume 21  
Number 7

# Partial Ordering




# Partial Ordering

Fig. 1.



# Partial Ordering

Fig. 1.




# Logical Clocks







# Logical Clocks




# Logical Clocks



# Logical Clocks




# Logical Clocks



# State Machines



# State Machines



# State Machines

| empty queue |



request(A, t0)

| A |



request(B, t1)

| A | B |



release(A, t3)

| B |

---

# State Machines

The problem of failure is a difficult one, and it is beyond the scope of this paper to discuss it in any detail. We will just observe that the entire concept of failure is only meaningful in the context of physical time. Without physical time, there is no way to distinguish a failed process from one which is just pausing between events. A user can tell that a system has “crashed” only because he has been waiting too long for a response. A method which works despite the failure of individual processes or communication lines is described in [3].

---



# Conclusion

